

(6) Basics of the C++ Programming Language

Nico Ludwig (@ersatzteilchen)

TOC

- (6) Basics of the C++ Programming Language
 - The Heap: Dynamic Memory and dynamic Array Allocation
 - Automatic versus Dynamic Arrays
 - A Glimpse of the Topic "Stack versus Heap"
 - "Geometric" Properties of the Heap and the Stack
 - Lost Pointers and Memory Leaks
 - Advanced Cstrings: Buffers, Concatenation and Formatting
- Sources:
 - Bruce Eckel, Thinking in C++ Vol I
 - Bjarne Stroustrup, The C++ Programming Language

Arrays have a Compile Time fixed Length

- We begin our discussion by remembering a limitation of arrays: they must have a length, which is fixed at compile time.
- Actually, we cannot define a C++ array with a length, which is set during run time, e.g. via user input

```
int count = 0;
std::cout<<"How many numbers do you want to enter?"<<std::endl;
std::cin>>count;
if (0 < count) {
    int numbers[count]; // Invalid (in C++)! The symbol count must be a compile-time constant!
    for (int i = 0; i < count; ++i) {
        std::cout<<"Enter number "<<(i + 1)<<": "<<std::endl;
        std::cin>>numbers[i];
    }
}
```

Good to know

Actually, some compilers, e.g. g++ accept this code, and the code also works. Those compilers have been extended to offer variable length arrays (VLAs), which can be created with a dynamic length. – But VLAs are no standard C++17 feature!

- More exactly, C++ does not support automatic arrays with a dynamic length!


The Correct Way of creating dynamic Arrays in Code

- Instead of using an automatic array, we have to create a dynamic array, where we can specify a dynamic length:

```
#include <cstdlib>

int count = 0;
std::cout<<"How many numbers do you want to enter?"<<std::endl;
std::cin>>count;
if (0 < count) { // Create a properly sized block in heap. The function std::malloc() returns a generic
    // pointer (void*) and we have to cast this generic pointer to the type we need.
    int* numbers = static_cast<int*>(std::malloc(sizeof(int) * count));
    if (numbers) { // Check, whether std::malloc() was successful.
        for (int i = 0; i < count; ++i) { // Loop over the dynamically created array:
            std::cout<<"Enter number "<<(i + 1)<<": ";
            // Use the block like an ordinary array, e.g. with the []-operator:
            std::cin>>numbers[i];
        }
        std::free(numbers); // When done with the array, it must be freed!
    }
}
```

```
Terminal
NicosMBP:src nico$ ./main
How many numbers do you want to enter?
4
Enter number 1:
1
Enter number 2:
2
Enter number 3:
3
Enter number 4:
4
NicosMBP:src nico$
```



- The core new thing in this code is the usage of the functions `std::malloc()` and `std::free()` from `<cstdlib>`:
 - `std::malloc()` takes over the creation of the array, accepting the length, which was specified by the user at run time.
 - `std::malloc()` returns a `void*` representing the created array, this generic pointer must be cast to an `int*` (to make it usable as `int[]`).
 - Because `numbers` is no longer an automatic array, we must remove it from memory manually with `std::free()`.
 - With automatic arrays, the memory consumed by the array was given back to the system automatically, when it left its scope.


4

- Consequently check the success of `std::malloc()`! – I.e. program defensively!
- Why do we need to cast here?
 - This is the first time we really need to cast. – Here we need to cast a pointer to memory of raw type to a pointer to memory of the type we need.
 - The `void*` represents a pointer to memory of unknown type; casting is required to get a type with which we can work. A `void*` is irrelevant by itself, we can't even dereference it. The only thing we can do with it is comparing it with other pointers or 0.
 - Objects of type `void*` are generally provided by the C++ (memory) system and must be "given back" to the system, after the programmer no longer needs it.
 - Interestingly the conversion from `void*` to another pointer type is seen as conversion between related types, so a `static_cast` is sufficient.

Graphical Representation of Stack- and Heap-Memory

- The arrays we have discussed up to now are so called automatic arrays.
 - Automatic arrays must have a length, which is defined at compile time!


```
const int length = 4;  
int numbers[length];  
numbers[0] = 1;  
numbers[1] = 2;  
numbers[2] = 3;  
numbers[3] = 4;
```



- The memory occupied by an automatic array resides on the stack, therefore, here we use blue boxes to depict them.

- The new way to create arrays, i.e. via `std::malloc()`, allows us to create arrays with a length defined at run time:

```
#include <cstdlib>  
  
int length = 4; // Assume value from run time ...  
int* numbers2 = static_cast<int*>(std::malloc(sizeof(int) * length));  
if (numbers2) {  
    numbers2[0] = 1;  
    numbers2[1] = 2;  
    numbers2[2] = 3;  
    numbers2[3] = 4;  
    std::free(numbers2);  
}
```



- The memory occupied by an array created dynamically with `std::malloc()` resides in the heap memory.
- Hence we use green boxes to depict portions of heap memory.
- Because this array is not automatically served back to the system, we have to call `std::free()` to remove it manually from memory.

5

- In this course we make a distinction between heap and free store. This distinction is not defined by the C++ standard. Instead it is used colloquially to tell the memory managed by `std::malloc()/std::free()` from that managed by `new/delete`, because they could work in different memory locations each.

std::malloc() in Focus

- In opposite to automatic memory, acquiring memory from the heap is a more explicit process.
 - The explicit acquiring of heap-memory is called memory allocation, hence the function named `std::malloc()` for memory allocation.


```
// <cstdlib>
namespace std {
    void* malloc(size_t size);
}
```

Good to know
from latin *allocare* or *locare*: "to place something"


- `std::malloc()`'s parameter size expects the size of the memory to allocate from the heap in byte (= `sizeof(char)`).
- However, the returned value is somewhat special: The returned value is of type `void*`. – Why that?
 - `void*` represents the generic pointer type. A `void*` can point to any kind of memory from any source (stack or heap).
 - `std::malloc()`'s function is to allocate a block of raw memory from the heap. Raw means that any type could be stored in this block.
 - In C++ we represent blocks of memory as arrays: an array is a list of elements, that reside in a contiguous block of memory.
 - The elements of this raw array have no type, they are just raw bytes (having `sizeof(char)`).
 - An array of elements without type, cannot decay to a pointer of a certain type, but can decay to the generic pointer-type `void*`.
- The important take away for the time being: dealing with heap memory means dealing with pointers.
- The other take away is, that this pointer is of type `void*`. – What can we do with it? How to "type" it?

std::malloc() and the generic Pointer Type (void*)

- All right, so we get a `void*` from `std::malloc()`. – How to progress from that?
- The first thing we have to do is telling the `void*`-referenced memory its type, currently, it is just a byte-array:



- We do this by casting the `void*` to a specific type:



- The cast to a specific pointer type tells `numbers`, that the referenced memory should be treated like an array of `int`!
- i.e. each element in `numbers` actually is an `int`, and has a size of 4(B).
- The cast works like a contact lens, which focuses the pointer from a blurry raw memory to an `int` array!
- Contact lens: Notice, that `memoryBlock` and `number` refer the same address, `numbers` just has a "sharper" view as `int` array.

Checking Heap-Memory Allocation and Freeing

- After `std::malloc()` returns a pointer, we have to check it for being a `nullptr`!
 - If `std::malloc()` didn't succeed, it'll return a `nullptr`.
 - `std::malloc()` fails, if there is no more heap memory for the requested allocation.
 - `std::malloc()` fails, if there is no contiguous portion in the heap memory to fit the requested block's in size.
 - (The `nullptr`-check can be done on the `void*` or the cast pointer "wearing the contact lens".)
- If `std::malloc()` succeeded, i.e. its result is not `nullptr`, we can use the result like any other array.
 - Esp. we can use the `[]`-operator to access/modify elements. I.e. the `[]`-operator can directly be used on a pointer (which is no `void*`).

```
int* numbers = static_cast<int*>(std::malloc(sizeof(int) * 4));
if (numbers) { // If the allocation succeeded ...
    // ...
}
```

```
if (numbers) { // If the allocation succeeded ... we can safely use numbers:
    numbers2[0] = 1;
    numbers2[1] = 2;
    numbers2[2] = 3;
    numbers2[3] = 4;
    // ...
}
```

- When we are done with the dynamic array `numbers`, we have to remove it from the heap with `std::free()`.
 - Because `std::free()` doesn't work with `nullptr`s, we must only call `free`, if `std::malloc()` succeeded!

```
if (numbers) { // If the allocation succeeded ... we can safely use numbers:
    // ...
    std::free(numbers); // free numbers from memory
}
```


8

- The `[]`-operator dereferences the pointer, but `void*` cannot be dereferenced, because the type to which it points is undefined!
- When we discussed `std::memcpy()`, we already had to deal with functions accepting `void*`s. But this time we have to work with returned `void*`s explicitly, which forces us to cast to other types, because `void*`s are pointing to raw memory. `std::malloc()`, `std::free()` and `std::memcpy()` are all low level functions.
- Can we assume that `std::malloc()` returns a pointer to a gap-free portion in the heap?
 - As far as our current knowledge is concerned: yes! It is required, because `std::malloc()` can be used to create arrays, and arrays need to represent contiguous blocks of memory.

Example: Why dynamic Memory is needed: Returning Arrays – Part I

- Of course the example we discussed up to now does not reflect a real life scenario.
- Consider following situation, in which we're going to return an array from a function.

```
int* getNumbers() {
    // Defining a function that returns a pointer to a
    int numbers[] = {1, 2, 3, 4}; // locally defined array (created on the stack).
    return numbers;               // This pointer points to the 1st item of values.
}
```


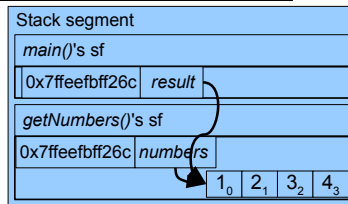


- When `getNumbers()` is called, it creates an `int[4]` on the stack.
- Then `getNumbers()` returns this array, it decays to `int*`, hence the return type.

- When `main()` calls `getNumbers()` it stores the returned `int*` in the local variable `result`.

```
int* result = getNumbers(); // Semantically wrong! result points to
std::cout<<"2. number is: "<<result[1]<<std::endl; // already discarded memory.
// The array "values" is gone away, result points to its scraps, probably rubbish!
```

- This won't work! When `getNumbers()` returns, its stackframe will be popped from the stack.
- When the sf is popped, all its locals will be removed from the stack as well, hence we call it autom. memory!
- `result` in `main()` will contain an invalid address pointing to unknown content!




Example: Why dynamic Memory is needed: Returning Arrays – Part II

- One way to make returning of array from a function possible is using dynamic memory.
- Let's rewrite `getNumbers()` to create the array in question in the dynamic memory.

```
int* getNumbers() {
    int* numbers = static_cast<int*>(std::malloc(sizeof(int) * 4));
    if (numbers) {
        numbers[0] = 1; numbers[1] = 2; numbers[2] = 3; numbers[3] = 4;
    }
    return numbers;
}
```

- When `getNumbers()` is called, it creates an `int[]` on the heap.





- When `main()` calls `getNumbers()` it stores the returned `int*` in the local `result`.

```
int* result = getNumbers();
if (result) {
    std::cout<<"2. number is: "<<result[1]<<std::endl;
    std::free(result);
}
```

- This works! When `getNumbers()` returns, its stackframe will be popped from the stack.

- But `getNumbers()`' return value is still a valid pointer!

- The locals in `getNumbers()` are popped from the stack, but the heap memory survived this!




- Once again: Consequently check the success of `std::malloc()`! – Program defensively!
- RAI means that a resource, which requires e.g. dynamic memory or other operating system resources, will be initialized and freed analogous to the lifetime of a variable. – Practically it means that a resource from dynamic memory can be controlled by a variable on the stack! – This can be implemented with user defined types.

Example: Why dynamic Memory is needed: Returning Arrays – Part III

- The graphical representation of the stack- and heap-situation shows that the heap is kind of shared among functions:

```
int* getNumbers() {  
    int* numbers = static_cast<int*>(std::malloc(sizeof(int) * 4));  
    if (numbers) {  
        numbers[0] = 1; numbers[1] = 2; numbers[2] = 3; numbers[3] = 4;  
    }  
    return numbers;  
}
```

```
int* result = getNumbers();  
if (result) {  
    std::cout<<"2. number is: "<<result[1]<<std::endl;  
    std::free(result);  
}
```




- Notice following points:
 - We have a pointer, that is a variable on the stack, but this time it holds the address to memory on the heap!
 - When we have an array from the heap, we can access/modify its elements with the []-operator, but other aspects have changed:
 - Using the heap comes for a price: since no stackframe controls the lifetime of memory, we have to control it manually!
 - I.e. we have to remove the memory we have allocated on the heap ourselves using explicit code, namely calling `std::free()`. 11
 - Also mind, that we have to check, whether allocation succeeded with a `nullptr`-check!

Wrap up: automatic and dynamic Arrays in Code


- Creation of automatic arrays:

```
int automaticArray[100];  
// This is a simple automatic array with a compile time constant size of 100.
```



- Creation and freeing of dynamic arrays:

```
int* dynamicArray = static_cast<int*>(std::malloc(sizeof(int) * 100));  
// - Indeed the syntax looks weird, not even similar to the autoArray example.  
// - The type of the variable we assign to is int-pointer.  
// - The function std::malloc() is used to create a raw memory-block in heap.  
// - std::malloc() returns a generic pointer (void*) to this raw memory-block, it does so, because  
//   it doesn't know, what the programmer wants to do.  
// - So, as programmers we need to tell C/C++ that we want to use the allocated memory  
//   block as int-array, therefor we need to cast the generic pointer (void*) to int-pointer. - We  
//   "put contact lenses on".  
if (dynamicArray) {  
    std::free(dynamicArray); // Free the dynamically created array in the right place.  
}
```



- Dynamically created arrays can have the length 0!
- Decay makes automatic arrays indistinguishable from dynamic arrays as function parameters!
 - E.g. we cannot blindly free a pointer passed to a function in the code of this function.
 - Specific implementation strategies and good documentation is required!


12

- Whether objects are stored on the heap or the stack is an implementation detail in many other languages, not so in C++, because the programmer explicitly decides, where an object will be created. (Annotated C# Language Reference, Anders Hejlsberg et. al.)

But, how does returning of "normal" Variables compute?

- Arrays can't be returned by value, so they can't be copied! – This is exactly the problem we solved with the heap!
- But taking a look back, how is the situation with non-arrays, i.e. "normal" automatic locals?
- When a local variable is returned from a function, it will be simply copied.

```
int getValue() { // Just returns an int:  
    int value = 42; // value is an automatic variable on the stack.  
    return value; // Returns value.  
}
```



- When `getValue()` ends, `value` will be popped from `getValue()'s` stack.
- But returning the content of `value` will push its content to the stackframe of the caller function. – This is the copy activity!
- => In effect `value` will be copied to its caller when `getValue()` returns.

```
void foo() { // Calls getValue():  
    int result = getValue(); // The returned int was pushed on foo()'s stack by getValue()  
                             // and will be copied into result.  
}
```

- When using automatic memory, there is no connection between the sfs, other than copying activities.
 - Esp. addresses are irrelevant, because `value` is getting copied and not shared somehow!

13

- Arrays can also be generated as **static** arrays instead of automatic arrays. A **static** array can be returned from a function. The lifetime of a **static** array is not restricted to a function's local scope.

Stack vs. Heap: It's not a Mystery, just two Concepts

- The stack is a conceptual place, where local (auto) variables reside.
 - This is a little oversimplification, but each function has its own portion of the stack, the so called stackframe.
 - The lifetime of a stack variable is visible by its scope (i.e. it is automatic: auto).
 - The stack is controlled by hardware and owned by hardware.
- The heap is a conceptual place, where all dynamic contents reside.
 - All functions of a program generally use the same heap.
 - Dynamic content must be created by the programmer manually.
 - The heap is controlled by software, the heap manager (`std::malloc()`, `std::free()` etc.).
 - There is always an "entity" that is in charge for the allocated heap memory.
 - This "entity" is responsible for explicit freeing the allocated heap memory.
 - In the end, the lifetime of a dynamic content is controlled by the entity's programmer.
- We'd try to control as less memory as possible manually: using the stack is preferred!

14

- What is a scope?
- See RAI!

Stack vs. Heap: In Memory (RAM)

- There is the illusion that all the machine's memory is owned by a program.
 - This is not true, each program uses its own portion of memory respectively.
 - But in the following graphics we stick to this illusion.
- The memory is segmented, different segments have different functions.
- Esp. the stack and heap segment often have special "geometric" properties:
 - The heap segment resides at lower addresses than the stack segment.
 - The addresses of subsequent stack variables are decreasing.
 - This is called "descending stack".
 - The stack evolves/grows to lower, the heap to greater addresses.
 - In fact, stack and heap grow to meet each other halfway!
 - Compared to the stack, the heap is very big: dynamic content is typically bigger than automatic content (e.g. local variables).

Stack vs. Heap: Conventional Locations in Memory



16

- Why $2^{32}-1$?
 - Well 2^{32} is 4,294,967,296 (ca. 4GB), but we have to subtract 1 in order to get space for the address 0!

The lost Pointer to the Heap Memory in Code

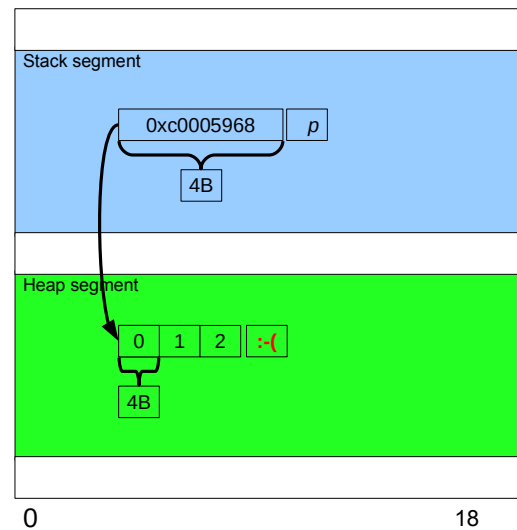
```
void f(int count) {  
    // Allocating an array of three ints. f() is in charge of the dynamic content, to which  
    // p points!  
    int* p = static_cast<int*>(std::malloc(sizeof(int) * count));  
    if (p) { // Check std::malloc()'s success.  
        for (int i = 0; i < count; ++i) {  
            p[i] = i;  
        }  
        // The auto variable p will go out of scope and will be popped from stack. But the  
        // referred dynamic content is still around!  
    }  
    //-----  
    // Calling f():  
    f(3);  
    // Oops, nobody did free the dynamic content, to which p was pointing to! Now there is  
    // no pointer to the dynamic content in avail. This is a semantic error, a memory leak of  
    // sizeof(int) * 3. The compiler will not see any problem here!
```

The lost Pointer to the Heap Memory in Memory

2³² - 1

```
void f(int count) { // Allocating an array of three ints.  
    int* p = static_cast<int*>(std::malloc(sizeof(int) * count));  
    if (p) { // Check std::malloc()'s success.  
        for (int i = 0; i < count; ++i) {  
            p[i] = i;  
        }  
    }  
}
```

```
// Calling f():  
f(3);  
// After f() did run: oops! The pointer to the allocated three  
// ints is lost, the allocated memory is orphaned. We have a  
// memory leak of 12B.
```



How to handle dynamic Content responsibly in Code

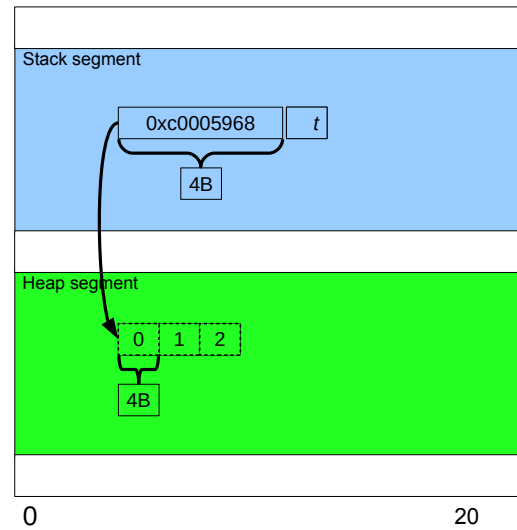
```
/**
 * Returns an int[] holding the first 'count' natural numbers.
 *
 * @param count the count of natural numbers to put into the result
 * @return a pointer to the int[] created on the heap, the caller must nullptr-check it
 * and std::free() it!
 */
int* g(int count) {
    // Allocating an array of three ints. g() is in charge of the dynamic content, to which
    // p points to!
    int* p = static_cast<int*>(std::malloc(sizeof(int) * count));
    if (p) { // Check std::malloc()'s success.
        for (int i = 0; i < count; ++i) {
            p[i] = i;
        }
    }
    // Returning p. Then g()'s caller is in charge of the dynamic content!
    return p; // The stack variable p will go out of scope and will be popped from the
    // stack. But the referred dynamic content is still around!
}
//-----
// Calling g():
int* t = g(3);
if (t) { // Fine! The returned pointer will be checked and freed correctly.
    std::free(t);
}
```

Handling dynamic Content responsibly in Memory

2³² - 1

```
/**
 * Returns an int[] holding the first 'count' natural numbers.
 *
 * @param count the count of natural numbers to put into the result
 * @return a pointer to the int[] created on the heap, the caller must nullptr-check it
 * and std::free() it!
 */
int* g(int count) { // Allocating an array of three ints.
    int* p = static_cast<int*>(std::malloc(sizeof(int) * count));
    if (p) { // Check std::malloc()'s success.
        for (int i = 0; i < count; ++i) {
            p[i] = i;
        }
    }
    return p; // This time: return p!
}
```

```
int* t = g(3); // Call g() and receive the pointer.
if (t) { // Check and free the content (i.e. the
    std::free(t); // memory from the heap).
}
// The local variable t is still on the stack.
```



Potential Problems with Heap Memory

- It is needed to check, whether allocation was successful!
- It is needed to free dynamic content in the right place manually.
 - We have to keep in mind that there is no garbage collection in C/C++.
 - So, we should not forget to free dynamically created content!
 - We should free dynamically created content as early as possible, but not too early!
 - We should not free dynamically created content more than once!
 - We should not free dynamically created content that we don't own.
- It's impossible to distinguish pointers to the stack from pointers to the heap.
 - Don't free pointers to the stack (i.e. pointers not from the heap)! -> It will result in undefined behavior!
- Secondary problem: we can't use array initializers when creating dynamic memory!
- Wherever functions deal with dynamically content, it must be documented where this memory must be freed!
 - Who's the owner of the memory?
 - Who's in charge of freeing the memory?

More Information about Heap Memory

- There exist two further useful functions to deal with heap memory (`<cstdlib>`):
 - `std::realloc()` resizes a given block of heap memory.
 - `std::calloc()` allocates a block of size * count and initiates all "items" with 0.
 - The returned value must be checked for `nullptr` and freed with `std::free()` respectively.
- Free store in C++:
 - In C++ the heap segment can also be used as free store.
 - The operators `new` and `delete` act as interface to the free store.
 - These operators represent C++' way to dynamically allocate/deallocate user defined types.
 - In general, C's heap memory and C++' free store are incompatible.
- Often 3rd party libraries invent own allocation/deallocation mechanisms:
 - to deal with platform specialities,
 - and/or to encapsulate usage of dynamic contents.

22

- `std::realloc()`:
 - Present items will be preserved up to the passed length.
 - Possibly a pointer to another location in the heap is returned, rendering the passed address invalid.
 - The returned pointer should be checked for 0, before it is assigned to the passed pointer variable. Don't do it like so: `p = std::realloc(p, 5000)`! If reallocation didn't succeed the memory to which the passed pointer `p` refers won't be touched.

Cstring Limitations: Assignment and Extension

- Assigning arrays to copy their content, esp. cstrings is not allowed in C++:

```
char aString[] = "Malta";  
const char aString2[] = "Teneriffa";  
aString = aString2; // Invalid! Array type 'char [6]' is not assignable
```

- What we want to have: copy the content of *aString2* to *aString*. – Mind, that arrays offer no copy semantics on assignment!
- Even if this would work at compile time, it wouldn't at run time, because *aString* has too few elements!
- (At least we defined *aString* being non-const, so that copying would at least be possible logically...)

- We cannot "extend" a cstring by appending or adding further cstrings to it:

```
char aString[] = "Weyland Yutani";  
aString += " at LV-426"; // Invalid! Invalid operands to binary expression ('char [15]' and 'const char [11]')
```

- Even if this would work at compile time, it wouldn't at run time, because *aString* has too few elements!
- (At least we defined *aString* being non-const, so that extending/appending would at least be possible logically...)

- In both cases show similar problems:

- We have to allocate memory on demand to create cstrings, which are larger than the original cstrings.
- The dilemma: we know the effective lengths of these cstrings only at run time.
- The solution: we must create this memory dynamically in the heap.

Dynamic Cstrings – Putting Heap Memory to work with Cstrings

- As cstrings are `char` arrays underneath, they share the limits of other arrays:
 - (1) We cannot `resize/extend` or `assign` cstrings.
 - (2) We cannot `return` an automatic cstring variable from a function.
- These limitations can be solved by usage of the heap memory:
 - To overcome limitation (1): Create a dynamically sized `char[]` on the heap to hold a modified/resized copy of the original cstring.
 - To overcome limitation (2): Copy a cstring into a dynamically sized `char[]` on the heap and `return` the pointer from a function.
 - Of course somebody must free the dynamically created memory, when it is no longer required.
- Sidebar: Peculiarities of cstrings, not directly shared with other arrays:
 - Cstrings are `0-terminated`, so the `very last char[]` item contains a `0`.
 - We can get a cstring's length (`std::strlen()`), this is impossible with other arrays.
 - As cstrings are `const char`-arrays, we can't modify them.
 - Indeed we can `return` a cstring literal from a function!

24

- Why is it possible to return a literal cstring from a function?
 - Cstring literals are stored in the `static` portion of memory (they are an example of a `static` arrays, which were mentioned earlier)! We'll discuss `static` memory in a future lecture.

Dynamic Cstrings – The Roles of `const char[]` and `char[]`

- When we need to create an array that must be filled afterwards, we cannot use arrays with `const` elements.
 - Instead we need arrays as modifiable buffers.
 - This is needed, cause we need to assign to the elements, in order to modify the content!
- So, cstrings are of type `const char[]`, their matching buffer type is `char[]`.
 - The buffers we allocate dynamically for `char`-based cstrings are always of type `char[]`.
- Where cstring functions accept `const char*`-params, we can safely pass `char[]` buffers; they will be decayed to `const char*`.
- To sum up (`char`-based cstrings):
 - Cstrings are of type `const char[]`, they're not modifiable
 - Cstring buffers are of type `char[]`, they're modifiable.

Dynamic Cstrings – Non-Dynamic Assignment/Copying

- For completeness, this code shows a naïve way to implement cstring assignment/copying with an automatic array:

```
const char* const initialString = "Malta";
const char* const aString2 = "Teneriffa";

char aString[256]; // A large buffer, with compile time fixed length.
std::strcpy(aString, initialString); // The "assignment"!
std::cout<<aString<<std::endl;

std::strcpy(aString, aString2); // The "assignment"!
std::cout<<aString<<std::endl;
```

```
Terminal
NicosMBP:src nico$ ./main
Malta
Teneriffa
NicosMBP:src nico$
```

- The code copies *initialString* and then *aString2* into *aString*.
- However, the problem is that *aString* has a length, which is fixed at compile time.
 - Problem 1: If a cstring we copy into *aString* is smaller than the capacity of *aString*, we waste memory!
 - Problem 2: If a cstring we copy into *aString* is larger than the capacity of *aString*, we overwrite memory and risk a buffer overrun!
- In most scenarios we don't know, how large the buffer must be to hold the effective result.
 - We could define a very large (automatic) buffer as shown above, but this is neither efficient nor safe.
 - The preferred way is to calculate the required buffer length exactly and create it dynamically!
 - On the following slides, we'll see some examples, where buffers must be created dynamically. – We begin with rewriting this code!

Dynamic Cstrings – Dynamic Assignment/Copying

- All right. Copying a cstring using the heap is a pretty simple case, but requires a lot of code:

```
const char* const initialString = "Malta";
const char* const aString2 = "Teneriffa";

char* aString = static_cast<char*>(std::malloc(sizeof(char) * (std::strlen(initialString) + 1)));
if (aString) {
    std::strcpy(aString, initialString); // The "assignment"!
    std::cout<<aString<<std::endl;
    std::free(aString);
}

aString = static_cast<char*>(std::malloc(sizeof(char) * (std::strlen(aString2) + 1)));
if (aString) {
    std::strcpy(aString, aString2); // The "assignment"!
    std::cout<<aString<<std::endl;
    std::free(aString);
}
```

Terminal

```
NicosMBP:src nico$ ./main
Malta
Teneriffa
NicosMBP:src nico$
```

- The code copies *initialString* and then *aString2* into *aString*, *aString* is dynamically sized to fit its new content respectively.
 - The code bloat makes the check of the correct allocation after *std::malloc()* and the need to free the allocated memory afterwards!
- The assignment/copying operation is a rather unimpressive call to *std::strcpy()*.
 - Notice: we could also have used *std::memcpy()* to do the copying, but this had required more size-calculation efforts in the code!

Dynamic Cstrings – Modified Copy

- A rich set of functions dealing with individual `chars` can be found in `<cctype>`.
- E.g. there are functions to change the casing of a `char` (these functions await an `int/char` and return an `int/char`):
 - `std::tolower()` and `std::toupper()`, their result must be cast to `char` for presentation.

```
char ch = 'X';  
// If ch is no upper case letter, the same char will be returned.  
std::cout<<ch<<" as lower case: "<<static_cast<char>(std::tolower(ch))<<std::endl;  
// >X as lower case: x
```

- With `std::toupper()` we can create an "upper-cased" variant of another cstring by copying and modifying that copy:

```
const char* oldText = "home"; // The original plain c-string.  
// Create a buffer in the heap, large enough to hold the original cstring. The buffer needs a size of:  
// sizeof(char) * (count of chars/letters + one byte for the 0-termination).  
char* newText = static_cast<char*>(std::malloc(sizeof(char) * (std::strlen(oldText) + 1)));  
if (newText) { // Check std::malloc()'s success.  
    // Loop over the original cstring and store the upper case variant of each char and the  
    // 0-termination into newText at the same index.  
    for (int i = 0; i < std::strlen(oldText) + 1; ++i) {  
        newText[i] = std::toupper(oldText[i]); // Modifies the non-const buffer.  
    }  
    std::cout<<"The modified text: "<<newText<<std::endl;  
    // >The modified text: HOME  
    std::free(newText); // Free the buffer.  
}
```

- Strategy: create a new buffer dynamically with the length of the original cstring + 1 (0-termination).
- Iterate over the buffer and the original cstring and set the upper-cased `char` of the original cstring at the new buffer's index.

28

- If we pass a non-letter `char` to `std::tolower()/std::toupper()` the passed `char` will just be returned.
- `std::tolower()/std::toupper()` return either the lower case or upper case variant of the passed character or the character itself, if there is no lower/upper variant. Therefore it is safe to pass the 0-termination, because it'll survive the function call and the 0-termination will be copied accordingly.
- Mind, that we can modify the copy of the cstring, because it resides in a modifiable/non-`const char`-buffer.

Dynamic Cstrings – String Concatenation – Part I

- Often it is required to dynamically compose cstrings by extending or concatenating other cstrings.

```
const char* s1 = "Weyland Yutani", *s2 = " at " *s3 = "LV-426";
```

- Concatenation (concat) means, that multiple cstrings are "jammed together" to build another cstring:

```
"Weyland Yutani" o " at " o "LV-426" = "Weyland Yutani at LV-426"
```

- Concat is not an addition, because a "real" mathematical/numerical addition is a commutative operation, string concat is surely not!

- The effective result we want to reach is to have a new cstring, which holds *s1*, *s2* and *s3* as concatenated cstring


- To make this work, we create a `char[]`-buffer, which is large enough to hold all chars of *s1*, *s2* and *s3* and the 0-termination!

- If we'd like to have this buffer exactly sized, we must use a dynamically allocated `char[]`!

```
// Calculate the length of the resulting cstring:
std::size_t countOfChars = std::strlen(s1) + std::strlen(s2) + std::strlen(s3);
// Allocate buffer with the exact size, sufficient for our situation (the lengths + 1):
char* buffer = static_cast<char*>(std::malloc(sizeof(char) * (countOfChars + 1)));
```

- Then, we have to concatenate the cstrings *s1*, *s2* and *s3* into the buffer with the function `std::strcat()` from `<cstring>`:

```
if (buffer) { // Check std::malloc()'s success then do the concatenation:
    buffer[0] = 0; // Set the 0-char at first char in the buffer.
    std::strcat(buffer, s1);
    std::strcat(buffer, s2);
    std::strcat(buffer, s3);
    std::cout<<buffer<<std::endl;
    // >Weyland Yutani at LV-426"
    std::free(buffer); // Free buffer.
}
```




Dynamic Cstrings – String Concatenation – Part II

- `char* strcat(char* destination, const char* source);` :
 - `std::strcat()` reads *destination* starting with the first `char`, until it finds the first 0-`char`.
 - Then it copies all `chars` from *source* into buffer until *source*'s 0-termination is found incl. *source*'s 0-termination.

```
char* buffer = static_cast<char*>(std::malloc(sizeof(char) * (countOfChars + 1)));
if (buffer) {
    buffer[0] = 0; // Set the 0-char at first char in the buffer.
    std::strcat(buffer, s1);
    std::strcat(buffer, s2);
    std::strcat(buffer, s3);
    std::free(buffer);
}
```

- Right after the *buffer* is created and *buffer[0]* set to 0, we have an empty array with a capacity of 25 elements.
 - After that follow up calls to `std::strcat()` fill the buffer in a concatenating manner:



30

- The memory representation is not precise on showing, how *buffer* changes. – Following the animation one could assume, that *buffer* is `std::free()`'d and created anew from the heap for multiple times, but instead it is modified in place by `std::strcpy()`!
- `std::strcat()` returns *destination*.

Dynamic Cstrings – Formatting – Part I

- Besides "simple" concatenation of cstrings we can also format cstrings dynamically.
 - Formatting is useful, if we want to guarantee a special and reusable look of a text with minimal effort.
 - Formatting is only about formatting of bare text data, it has nothing to do with font styles like "bold" or "underlined"!
- To format cstrings, we will use the function `std::sprintf()` from `<cstdio>`.
 - `std::sprintf()` awaits a buffer into which the formatted result will be stored, a format string and the values to be formatted as cstring:

```
char fullText[256];  
std::sprintf(fullText, "Duration: %dms %s", 3 * 4, "speed-test");  
// fullText = "Duration: 12ms speed-test"
```

- The format string works like a template, which defines the result cstring with placeholders, which are replaced by values.
 - The placeholders in the format string template are denoted by the %-character and a so called format specifier.
 - The placeholders are effective to the values in the order they are written: `%d` to `3 * 4` and `%s` to `"speed-test"`:

```
std::sprintf(fullText, "Duration: %dms %s", 3 * 4, "speed-test");
```

- Here we have two placeholders `%d` and `%s`, which address the arguments `3 * 4` and `testName` respectively.
 - `%d` tells `std::sprintf()` to handle its corresponding value (`3 * 4`) as decimal integer, `%s` denotes a cstring (`"speed-test"`).
- `std::sprintf()` can principally deal with an unlimited amount of placeholders and can replace an unlimited amount of values.
 - But, this would mean that `std::sprintf()` deals with an unlimited amount of argument, how can that work?

- String-formatting: integrals and floats are promoted to int/double, when passed to ..., therefor `%d` and `%g` can accept either.

Dynamic Cstrings – Formatting – Part II

```
std::sprintf(fullText, "Duration: %dms %s", 3 * 4, "speed-test");
```

- Here we call `std::sprintf()` with 4 args, which can be split into buffer and format as first 2 args and a list of values.
 - The signature of `std::sprintf()` looks like this:

```
int sprintf(char* buffer, const char* format, ...);  
  
char fullText[256];  
std::sprintf(fullText, "Duration: %dms %s", 3 * 4, "speed-test");
```

- As can be seen the parameter format stores the format string and the variable arity parameter args stores the remaining arguments.

- `std::sprintf()` can be used as alternative to `std::strcat()` to concat cstrings with less code:

```
const char* s1 = "Weyland Yutani", *s2 = " at ", *s3 = "LV-426";  
std::size_t countOfChars = std::strlen(s1) + std::strlen(s2) + std::strlen(s3);  
char* buffer = static_cast<char*>(std::malloc(sizeof(char) * (countOfChars + 1)));  
  
if (buffer) {  
    buffer[0] = 0;  
    std::strcat(buffer, s1);  
    std::strcat(buffer, s2);  
    std::strcat(buffer, s3);  
    std::cout<<buffer<<std::endl;  
    // >Weyland Yutani at LV-426"  
    std::free(buffer);  
}
```

```
const char* s1 = "Weyland Yutani", *s2 = " at ", *s3 = "LV-426";  
std::size_t countOfChars = std::strlen(s1) + std::strlen(s2) + std::strlen(s3);  
char* buffer = static_cast<char*>(std::malloc(sizeof(char) * (countOfChars + 1)));  
  
if (buffer) {  
    std::sprintf(buffer, "%s%s%s", s1, s2, s3);  
    std::cout<<buffer<<std::endl;  
    // >Weyland Yutani at LV-426"  
    std::free(buffer);  
}
```


Excursus: Variable Length Argument Lists in C

- Esp. C is well known for its feature of functions, which can cope with variable argument lists (vargs):

```
#include <stdarg.h>
// C variadic function example
int sum(int nNumbers, ...) {
    int sum = 0;
    va_list args;
    va_start(args, nNumbers);
    for (int i = 0; i < nNumbers; ++i) {
        nSum += va_arg(args, int);
    }
    va_end(args);
    return sum;
}

const int full_sum = sum(3, 1, 2, 3);
// full_sum = 6
```

Good to know

All standard C/C++ functions have the calling convention `__cdecl`. Only `__cdecl` allows variable argument lists, because only the caller knows the argument list and only the caller can then pop the arguments. `__stdcall` functions execute a little bit faster than `__cdecl` functions, because the stack needs not to be cleaned on the callee's side (i.e. within a `__stdcall` function).

- Featured by the ubiquitous function `std::sprintf()`, applied via the `...`-operator (ellipsis-operator).
 - The mandatory vargs' first argument must be interpreted, to guess how many vargs follow.
 - Vargs are harmful: It's a way to introduce security leaks through stack overruns. (Just call `sum(4, 1, 2, 3)` and see what happens.)
- How does it work? The vargs-features works very near the metal:
 - The compiler calculates the required stack depending on the arguments and decrements the stack pointer by the required offset.
 - As arguments are laid down on the stack from right to left, `nNumbers` is on offset 0.
 - Then `nNumbers` is analyzed and the awaited offsets are read from the stack. Here an offset of, e.g., 4B for each `int` passed to `sum()`.

33

- When we think about `std::sprintf()`, we don't pass the count of arguments via an explicit argument, instead `std::sprintf()` analyzes the passed format-string and determines the amount of placeholders.
- The calling convention `__cdecl` is a C/C++ compiler's default, `__stdcall` is the calling convention of the Win32 API, because it works better with non-C/C++ languages. `__cdecl` requires to prefix a function's name with an underscore when calling it (this is the exported name, on which the linker operates). A function compiled with `__stdcall` carries the size of its parameters in its name (this is also the exported name). – Need to encode the size of bytes or the parameters: If a `__cdecl` function calls a `__stdcall` function, the `__stdcall` function would clean the stack and after the `__stdcall` function returns the `__cdecl` function would clean the stack again. – The naming of the exported symbol of `__stdcall` functions allow the caller to know how many bytes to "hop", because they've already been removed by the `__stdcall` function. Carrying the size in a function name is not required with `__cdecl`, because the caller needs to clean the stack. – This feature allowed C to handle variadic functions with `__cdecl` (nowadays the platform independent variadic macros can be used in C and C++).
- Other calling conventions:
 - pascal: This calling convention copies the arguments to the stack from left to right, the callee needs to clean the stack.
 - fastcall: This calling convention combines `__cdecl` with the usage of registers to pass parameters to get better performance. It is often used for `inline` functions. The callee needs to clean the stack. The register calling convention is often the default for 64b CPUs.
 - thiscall: This calling convention is used for member functions. It combines `__cdecl` with passing a pointer to the member's instance as if it was the leftmost parameter.
- In this example the RV (EAX on x86) register can only store values of 4B. In reality the operation can be more difficult.
 - For floaty results the FPU's stack (ST0) is used.
 - User defined types (e.g. `structs`) are stored to an address that is passed to the function silently.
- It is usually completely different on micro controllers.

Formatting Cstring – more Examples

- For string formatting we can use a lot of format specifiers different from %s or %d, esp. for formatting floaty values.

- The general format for floaty values, such as `double` is specified with the `%g` placeholder as format specifier:

```
char buffer[256];
double value = 2234567890;
std::sprintf(buffer, "value: %g", value);
// buffer = "2.23457e+09"
```

- The default precision, which is effective on all digits is 6. We can change this to a precision of 10, with the `%.10g` format specifier:

```
std::sprintf(buffer, "value: %.10g", value);
// buffer = "2234567890"
```

- With the format specifier `%e` we specify the scientific format for the corresponding floaty value:

```
double value2 = 2.23456789;
std::sprintf(buffer, "value2: %e", value2);
// buffer = "2.234568e+00"
```

```
// Scientific notation with precision of 2:
std::sprintf(buffer, "value2: %.2e", value2);
// buffer = "2.23e+00"
```

- With the format specifier `%f` we specify the fixed point format for the corresponding floaty value:

```
std::sprintf(buffer, "value2: %f", value2);
// buffer = "2.234568"
```

```
// Floating point value rounded to two digits:
std::sprintf(buffer, "value2: %.2f", value2);
// buffer = "2.23"
```

34

- `std::sprintf()`:
 - Superfluous arguments in the varargs will be ignored.
 - Too few arguments result in undefined behavior.
 - If the format-string and the arguments don't match the behavior is undefined.
 - There are different ways to control the format in a more detailed fashion. Besides the precision, additional flags control alignment and padding, also the width and the length can be controlled. – But the resulting format strings quickly become difficult to read.
 - `%p` can be used to output pointers.
 - Downsides:
 - It should be said, that `std::sprintf()` must also be used with care, because one can easily overwrite the passed buffer by getting format specifiers wrong. – It can be tricky to calculate the correct size of the buffer.
 - Currently, C++ does not allow to "invent" new format specifiers of own UDTs. However, this issue was addressed by the possibility to overload `operator-<<` for own UDTs.

Thank you!